



Image Processing Vision System for Manned and Unmanned Aircraft

CS 463 - Senior Capstone Final Report

June 10th, 2016

Group 4 - HawkEye Crew

Hailey Palmiter

Scott Griffy

Ryan Kitchen

Abstract

This document is a final report constructed using the IEEEtran style guidelines. The report provides an introduction to our senior project, a project description, and a personal overview from each member of the team about what they learned over the course of the project. The report includes our original requirements document, design document, and technological review with an added discussion about each document. This discussion describes how the project had changed over the course of the year in comparison to how it was perceived during the planning process. Also included is our weekly blog posts that had been generated over the course of the year to keep our TA and our sponsors up to date with our progress, and our final poster that was used at the Engineering Expo. Appendix 1 covers essential code listings and Appendix 2 displays the hardware we were able to interact with and a few output images from our vision system.

CONTENTS

I	Introduction	4
II	Original Requirements Document	4
II-A	Introduction	4
II-A1	Purpose	4
II-A2	Scope	4
II-A3	Definitions	5
II-A4	References	5
II-A5	Assumptions and Constraints	5
II-B	Overall Description	5
II-B1	Product Perspective	5
II-B2	Product Functions	5
II-B3	User Characteristics	5
II-B4	Constraints	6
II-B5	Assumptions and Dependencies	6
II-B6	Stretch Goal Timeline	6
II-C	Specific Requirements	6
II-C1	External Interfaces	6
II-C2	Functions	6
II-C3	Performance Requirements	6
II-C4	Design Constraints	6
II-C5	Maintainability	6
II-C6	Portability	6
II-D	Stretch Goals	7
II-E	Gantt Chart	7
III	Requirements Document Changes	7
III-A	Final Gantt Chart	7
IV	Design Document	13
IV-A	Introduction	13
IV-A1	Scope	13
IV-A2	Purpose	13
IV-A3	Intended Audience	13
IV-A4	References	13
IV-B	Definitions	13
IV-C	Conceptual Model for Software Description	14
IV-C1	Software Design in Context	14

	IV-C2	Software Design Descriptions within the Life Cycle	14
IV-D		Design Architecture	14
	IV-D1	Stakeholder Concerns and Requirements	14
	IV-D2	Description of Architectural Design	14
	IV-D3	Validation of Design	14
	IV-D4	Overview of Viewpoints and Design Languages	15
IV-E		Design Viewpoints	15
	IV-E1	Context Viewpoint	15
	IV-E2	Structure Viewpoint	16
	IV-E3	Interaction Viewpoint	18
	IV-E4	Information Viewpoint	19
	IV-E5	State Dynamics Viewpoint	22
V		Design Document Changes	23
V-A		Overall Description Changes	23
V-B		Specific Requirements	23
V-C		Information Viewpoint	23
VI		Technology Review	23
VI-A		Camera Interface (Hardware Level)	23
	VI-A1	USB 3.0	24
	VI-A2	Gigabit Ethernet	24
	VI-A3	MIPI Camera Interface	24
	VI-A4	Decision	25
VI-B		Video Subsystem (Software Level)	25
	VI-B1	OpenCV	25
	VI-B2	Direct Memory Handling	25
	VI-B3	Gstreamer	25
	VI-B4	Decision	26
VI-C		Image Processing	26
	VI-C1	On-board GPU	26
	VI-C2	Integrate Stream Processor	26
	VI-C3	PCIe GPU	27
	VI-C4	Decision	27
VI-D		Demo Interface	27
	VI-D1	QT Framework	27
	VI-D2	OpenCV GUI	27
	VI-D3	OpenGL	28
	VI-D4	Decision	28

VII	Technology Review Document Changes	28
VII-A	Technology Changes	28
VII-A1	Camera Interface	28
VII-A2	Video Subsystem	28
VII-A3	Image Processing	28
VII-A4	Demo Interface	28
VIII	Weekly Blog Posts	29
VIII-A	Fall Term 2015	29
VIII-A1	October 16, 2105	29
VIII-A2	October 23, 2105	29
VIII-A3	October 30, 2105	29
VIII-A4	November 6, 2105	29
VIII-A5	November 13, 2105	30
VIII-A6	November 20, 2105	30
VIII-A7	November 27, 2105	31
VIII-A8	December 4, 2105	31
VIII-B	Winter Term 2016	31
VIII-B1	January 8, 2016	31
VIII-B2	January 15, 2016	32
VIII-B3	January 22, 2016	33
VIII-B4	January 29, 2016	34
VIII-B5	February 5, 2016	34
VIII-B6	February 12, 2016	35
VIII-B7	February 19, 2016	35
VIII-B8	February 26, 2016	35
VIII-B9	March 4, 2016	36
VIII-B10	March 11, 2016	36
VIII-C	Spring Term 2016	36
IX	Final Poster	38
X	Project Documentation	39
XI	Discovery of New Technologies	39
XII	Personal Experiences	39
XII-A	Hailey Palmiter	39
XII-B	Scott Gruffy	39
XII-C	Ryan Kitchen	39

Rockwell Collins, our project sponsor, designs video vision systems for pilots to use during flight. Pilots often use these enhanced image systems to help them see better in rough weather conditions, and to generally assist during different flight operations, such as landing. An example of this functionality could be overlaying a pilot's view with graphics to help the pilot locate a landing strip in a storm, or automatically turning a night-vision camera off when it is not needed to save power. The hardware used for this system must limit the power consumption it pulls from the aircraft, and also must be lightweight in order to have the least affect on the airplane. Even an extra five pounds added to the plane during a year can add up to thousands of dollars in fuel costs. In order to provide pilots with this specific low power, low weight, and enhanced image processing, Rockwell Collins develops software on Field Programmable Gate Arrays (FPGAs). This piece of hardware makes the code very complex and very costly to develop. New vision enhancements can take weeks or months to develop on the FPGAs. The FPGAs are currently Rockwell Collins only option that meets the requirements needed to create practical systems that pilots can use effectively.

Our goal is to provide a proof of concept for an alternative to Rockwell Collins' FPGAs. It must meet the performance metrics, provide faster implementation time, and reduce the cost of production. Specifically, we are designing a proof of concept using single board computers (SBCs). Single board computers differ from FPGAs because they have a standardized execution environment, which allows simplified code to be executed on it, reducing the development time. Single board computers also use low cost hardware and don't consume much power. If we can prove that the video quality produced by the SBCs is adequate for pilots to use, they will meet all the requirements needed to develop a practical vision system on. In the air, FPGAs often take feeds from multiple cameras and run a lot of processing algorithms on those images. In order to be effective, the SBCs should also be able to handle this operation. Our project aims to test and measure the capabilities of a single board computer by delivering a multiple-stream video display that has been processed to a high degree. We believe the best candidate for this is to use NVIDIA's single board computer, the Jetson TX1. Our goal is to fully test the Jetson TX1's ability to provide enhanced imaging. This proof of concept will result in measurements that will help Rockwell Collins determine the practicality of using single board computers for their vision systems.

II. ORIGINAL REQUIREMENTS DOCUMENT

A. Introduction

1) *Purpose:* Rockwell Collins currently develops their vision systems on Field Programmable Gate Arrays, making the time from design to product very long. Rockwell Collins is considering using Single Board Computers as a way to speed up this cycle. Our goal is to provide a proof of concept for a vision system on a Single Board Computer for Rockwell Collins. This proof of concept should help Rockwell Collins measure the practicality of Single Board Computers for vision systems.

2) *Scope:* Development will be limited to the Jetson TK1 or Jetson TX1. Evaluation of hardware platforms other than the TK1 and TX1 will be a stretch goal. In which case, these boards will serve as the base unit for comparison.

Software shall interface to a minimum of 2 cameras, receive video, and display it on a monitor while logging metrics. Metrics will be frames per second and operations per second.

3) *Definitions:*

SBC Single board computer
RC Rockwell Collins
FPS Frames per second
FPGA Field-programmable gate array
USB Universal serial bus
PoC Proof of concept
SVS Simple vision system

4) *References:* Carlo Tiana, Airborne Vision Systems Expert at Rockwell Collins Weston Lahr, Senior Software Engineer at Rockwell Collins, Jetson TK1 Embedded Hardware, NVIDIA

5) *Assumptions and Constraints:* We're assuming that either Jetson model can handle some sort of camera vision system. From initial investigations it appears that the Jetson TX1 is a likely candidate to support our project goals.

B. *Overall Description*

This section will give a bird's eye view of the project, outlining the major features and their purpose.

1) *Product Perspective:*

- a) *System Interfaces:* Linux, Jetson Distro
- b) *User Interfaces:* While using our deliverable, the user should be able to view output image from the Jetson.
- c) *Software Interfaces:* Our deliverable should be a demonstrate of the Jetson's video output, charts that display capability performance, information about the complexity and amount of video processing operations. A stretch goal would include to have the same output of the software running on a second screen.
- d) *Hardware Interfaces:* We must choose an interface between our program and the chosen camera attached to the Jetson SBC. We must also interface with the relative camera processing systems on the Jetson.
- e) *Communication Interfaces:* We will communicate with at least two different camera streams, from two separate cameras attached to our SBC.

2) *Product Functions:* The delivered demonstration should accurately demonstrate the performance limitations of the Jetson TK1 or TX1. Should the timeframe allow, we are to continue investigating other SBCs and providing a similar overview of other researched SBCs.

3) *User Characteristics:* Our users will be RC developers, looking to implement a given filter for vision processing platforms.

4) *Constraints*: Our deliverable must run on either the Jetson TK1 or Jetson TX1, and should simply test the Jetson's capabilities. We will use 1080p cameras.

5) *Assumptions and Dependencies*: We are assuming that the Jetson can at least handle some level of video processing. We are also going to rely on libraries and the operating system provided by Nvidia for the Jetson.

6) *Stretch Goal Timeline*: Stretch goals defined throughout this document may be delayed until all other requirements are finished and are not required to be in the deliverable at any point.

C. Specific Requirements

This section will outline some of the more specific parts of the project along with exact requirements that the deliverable should meet.

1) External Interfaces:

Two 1080p Cameras These will be used to transfer video data to the Jetson to be processed and displayed on a monitor.

Monitor Any monitor able to adequately display the output of the chosen SBC for the purposes of evaluation.

Jetson Should be able to read and process video data and stream it out to a monitor.

2) *Functions*: The system shall take input from the cameras and perform as much video processing on it as possible before dipping under the 30 fps minimum. If the system can't handle 30fps without any extra video processing, it should simply be displayed at the fastest fps possible.

3) *Performance Requirements*: The system should support at least 2 cameras with a stretch goal of 3 cameras. The system should support video output on one monitor, and if possible the video should be displayed at 30fps.

4) *Design Constraints*: Since this is a proof of concept, the only real design constraint is that it should be able to correctly assess and demonstrate the hardware limitations of the Jetson SBC.

5) *Maintainability*: Maintainability, outside of our timeframe, is a low-priority requirement as we are only delivering a proof of concept for the capabilities of our SBC.

6) *Portability*: Our software should give some idea of how similar code would perform on other platforms, even though the PoC itself doesn't have to be portable. The same goes for the camera and monitor. In order to help our stretch goal of evaluating a second SBC, it would help if the code, cameras, and monitor were all collectively portable to the second platform.















D. Stretch Goals

Some stretch goals were listed throughout the requirements. One purpose of the stretch goals is to further increase the effectiveness of the proof of concept. These added goals further stress the limitations of the Jetson TK1 or TX1, to further expand the capabilities of these SBC. A few other stretch goals not mention above would be implementing the same system on a similar SBC, and comparing that with the Jetson TK1 or TX1, or implementing object tracking as a video processing operation to test performance of the SBCs.

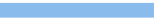


















E. Gantt Chart














III. REQUIREMENTS DOCUMENT CHANGES

A. Final Gantt Chart

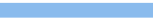


















ID	 Task Mode	Task Name	Duration	Start	Finish	Qtr 4, 2015 Oct	Nov	Dec	Qtr 1, 2016 Jan	Feb	Mar	Qtr 2, 2016 Apr
1		Start developing on the Jetson Tk1 board	49 days	Sun 11/1/15	Wed 1/6/16							
2		Install the libraries/software necessary for development	3 days	Mon 1/4/16	Wed 1/6/16							
3		Determine what libraries would be useful for a SVS	2 days	Mon 1/4/16	Tue 1/5/16							
4		Determine the environment necessary for these libraries	2 days	Mon 1/4/16	Tue 1/5/16							
5		Ensure that these libraries will allow us to create each part of the project	2 days	Tue 1/5/16	Wed 1/6/16							
6		Create a communal space for development	22 days	Sun 11/1/15	Mon 11/30/15							
7		Give everyone access to this space	20 days	Sun 11/1/15	Thu 11/26/15							


















Project: yah
Date: Mon 6/6/16

Task		Inactive Summary		External Tasks	
Split		Manual Task		External Milestone	
Milestone		Duration-only		Deadline	
Summary		Manual Summary Rollup		Progress	
Project Summary		Manual Summary		Manual Progress	
Inactive Task		Start-only			
Inactive Milestone		Finish-only			

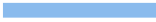


















ID	 Task Mode	Task Name	Duration	Start	Finish	Qtr 4, 2015 Oct	Nov	Dec	Qtr 1, 2016 Jan	Feb	Mar	Qtr 2, 2016 Apr
8		Create a plan for hardware, using a shared computer or individual computers	22 days	Sun 11/1/15	Sat 11/28/15							
9		Implement this plan, getting a shared computer or installing the necessary software on each of our computers	22 days	Sun 11/1/15	Mon 11/30/15							
10		Develop the camera vision system	25 days	Mon 1/4/16	Fri 2/5/16							
11		Get a working stream of video on the chosen platform	25 days	Mon 1/4/16	Fri 2/5/16							
12		Connect cameras to single board computer	13 days	Mon 1/4/16	Wed 1/20/16							
13		Display camera feed on output monitor	25 days	Mon 1/4/16	Fri 2/5/16							






















Project: yah
Date: Mon 6/6/16

Task		Inactive Summary		External Tasks	
Split		Manual Task		External Milestone	
Milestone		Duration-only		Deadline	
Summary		Manual Summary Rollup		Progress	
Project Summary		Manual Summary		Manual Progress	
Inactive Task		Start-only			
Inactive Milestone		Finish-only			




















ID	 Task Mode	Task Name	Duration	Start	Finish	Qtr 4, 2015 Oct	Nov	Dec	Qtr 1, 2016 Jan	Feb	Mar	Qtr 2, 2016 Apr
14		Develop method for measuring overall latency	25 days	Mon 1/4/16	Fri 2/5/16							
15		Measure fixed latency of camera and display	19 days	Mon 1/4/16	Thu 1/28/16							
16		Implement input-to-output timestamp of each frame	25 days	Mon 1/4/16	Fri 2/5/16							
17		Develop the benchmarking and testing user interface	31 days	Mon 1/4/16	Mon 2/15/16							
18		Display latency, frames per second, and CPU/GPU usage on screen	28 days	Mon 1/4/16	Wed 2/10/16							
19		Log with timestamps	26 days	Mon 1/4/16	Sat 2/6/16							
20		Display in real time	26 days	Mon 1/4/16	Mon 2/8/16							
21		Implement video combination algorithm	31 days	Mon 1/4/16	Mon 2/15/16							

Project: yah
Date: Mon 6/6/16

Task		Inactive Summary		External Tasks	
Split		Manual Task		External Milestone	
Milestone		Duration-only		Deadline	
Summary		Manual Summary Rollup		Progress	
Project Summary		Manual Summary		Manual Progress	
Inactive Task		Start-only			
Inactive Milestone		Finish-only			

ID		Task Mode	Task Name	Duration	Start	Finish	Qtr 4, 2015 Oct	Nov	Dec	Qtr 1, 2016 Jan	Feb	Mar	Qtr 2, 2016 Apr
22			Implement color filter layers	31 days	Mon 1/4/16	Mon 2/15/16							
23			Implement video postprocessing algorithms	4 days	Wed 2/10/16	Sat 2/13/16							
24			Blur removal	4 days	Wed 2/10/16	Sat 2/13/16							
25			Erosion	4 days	Wed 2/10/16	Sat 2/13/16							
26			Smoothing	4 days	Wed 2/10/16	Sat 2/13/16							
27			And more!	4 days	Wed 2/10/16	Sat 2/13/16							
28			Implement object tracking for a specified target (optional)	2 days	Sat 2/13/16	Mon 2/15/16							
29			Detect object with cameras 1 and/or 2	2 days	Sat 2/13/16	Mon 2/15/16							
30			Use camera with telescopic lens to zoom in and track target	2 days	Sat 2/13/16	Mon 2/15/16							
31			Benchmarking and evaluation	32 days	Sat 2/13/16	Mon 3/28/16							

Project: yah
Date: Mon 6/6/16

Task		Inactive Summary		External Tasks	
Split		Manual Task		External Milestone	
Milestone		Duration-only		Deadline	
Summary		Manual Summary Rollup		Progress	
Project Summary		Manual Summary		Manual Progress	
Inactive Task		Start-only			
Inactive Milestone		Finish-only			

IV. DESIGN DOCUMENT

A. Introduction

1) *Scope:* This software will implement a framework to develop, test, and benchmark video processing algorithms. The users will be able to test different combinations of algorithms using one or multiple camera inputs and produce output to one or multiple display windows. This software will be used by HawkEye Crew to determine if the Nvidia Jetson TK1 or other off the shelf single board computers have high enough performance to replace FPGA-based video processing systems currently used by Rockwell Collins.

2) *Purpose:* This software description document will provide HawkEye Crew with a road map to complete development of the software and fulfill the requirements in our Software Requirements Specification. This document explains how the system is going to work, who is going to use it, and how it is meant to be used.

3) *Intended Audience:* The intended audience of this design document is the developers who will design the system *us* and the sponsors of this project at Rockwell Collins.

4) *References:* IEEE. IEEE Std 1016-2009 IEEE Standard for Information Technology System Design Software Design Descriptions. IEEE Computer Society, 2009

B. Definitions

SBC Single board computer

RC Rockwell Collins

FPS Frames per second

FPGA Field-programmable gate array

USB Universal serial bus

PoC Proof of concept

UML Unified Modeling Language

DFD Data Flow Diagram: Shows how data moves between different components in the system

ER Diagram Entity Relationship Diagram: Shows how different data structures within the software are connected to each other

SDD Software Design Description

SRS Software Requirements Specification

SVS Simple Vision System

YUV A color space typically used as part of a color image pipeline

JSON JavaScript Object Notation

Modular Video Processing System The HawkEye Video Processing System's internal video processing algorithm management and execution system.

C. Conceptual Model for Software Description

1) *Software Design in Context:* For the HawkEye Video Processing System, we will be using a functional design method. Our system is designed to be modular and have standardized interfaces, so we will be able to develop new features and drop them in without modifying existing code. The reasoning for using a functional design rather than object oriented is because our software is focused on processing individual video frames in real time. Each algorithm instance will run once and has no state, only input and output data. For that reason functional programming is the best choice for this design.

2) *Software Design Descriptions within the Life Cycle:* This document details how we are going to implement software to meet our requirements. During the course of development, our requirements and our SRS may be updated. This document includes support for both our current requirements and provides room to add additional requirements. It also adds additional operational requirements and functionality beyond those specified in the requirements document. This document will also provide a reference for the creation of our testing plan.

D. Design Architecture

1) *Stakeholder Concerns and Requirements:* Our client Rockwell Collins would like to know if single board computers are a viable option to implement simple vision systems on. They've enlisted us, the HawkEyeCrew, to implement a proof of concept to measure the effectiveness of single board computers in this context.

We, the HawkEye Crew, will be using this software to test various implementations of video processing algorithms. In order for our system to be a valid proof of concept, our system will need to meet certain performance benchmarks and functional requirements described by our SRS. Our design concerns will be meeting these requirements.

2) *Description of Architectural Design:* Our system is designed to be a framework and testbed for the implementation of video processing algorithms. This design handles the hardware input interface, video output to the monitor, data flow between different algorithms, parallel operation of multiple operations, and benchmarking of individual components and overall aggregate timing.

In order to provide maximum flexibility and testing capabilities, the video processing algorithms and input/output devices are built as separate, interchangeable modules. These different modules are organized by the user in a configuration file, which determines both how the modules connect to each other and how the log file is formatted. This enables us to track the performance of individual algorithms and maximize our CPU and GPU usage while maintaining the performance requirements for this project.

3) *Validation of Design:* The core requirements of our system that this design either implements or implements the capability of testing are as follows:

- 1) System must be capable of processing multiple video streams
- 2) System must enable thorough testing of the Jeston's video processing capabilities
- 3) Frame rate must be at least 30FPS

- 4) End to end latency must be less than 100ms

These requirements are met by the following architectural features:

- 1) The modular video processing system enables multiple input and output devices.
- 2) The modular video processing system will enable us to maximize the usage of CPU and GPU processing power, which will show the maximum performance of the Jetson.
- 3) This software design includes built in speed benchmarking, so we will be able to track FPS.
- 4) This benchmarking system also produces latency details, so we will be able to determine if this system is capable of meeting the latency requirements Rockwell Collins is looking for.

4) *Overview of Viewpoints and Design Languages:* These viewpoints have been chosen to provide a complete description of the design and show how the design is compliant with the SRS. Each one provides details crucial to understanding how the design works and is a reference for implementation.

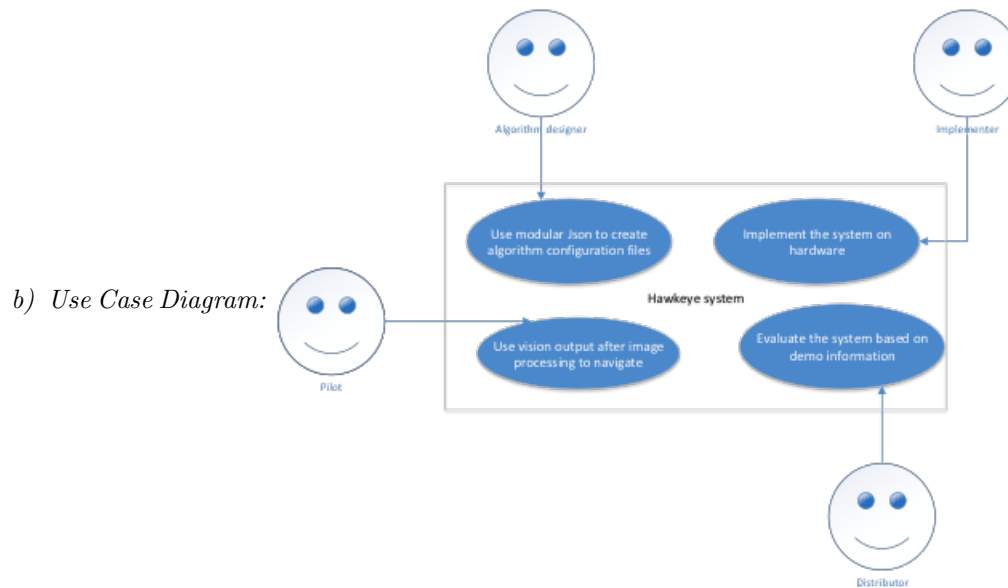
- 1) **Context viewpoint:** This viewpoint shows the different potential users of the software and how they would interact with our system.
Design Languages: Use Case Diagram
- 2) **Structure viewpoint:** This viewpoint shows how the streaming video flows through the system, and identifies the internal and external data connection points in the system. This viewpoint also shows the components of the system which enable connection of cameras via USB as per the SRS.
Design Languages: Data Flow Diagram
- 3) **Interaction viewpoint:** This viewpoint shows the order of operations on processing a video frame, as well as how the timing is integrated to meet the performance benchmarking requirement.
Design Languages: UML Sequence Diagram
- 4) **Information viewpoint:** This viewpoint details how the modular video processing system determines the order of operations for the various elements and algorithms it can create. It shows how the system meets the requirement for modularity and the capability to use multiple input and output devices.
Design Languages: Entity Relationship Diagram
- 5) **State Dynamics Viewpoint:** While the algorithm implementations in our design are stateless, the system itself is not. This viewpoint shows the transitions between different states and provides a road map for different states which will need to be individually tested.
Design Languages: UML State Transition Diagram

E. Design Viewpoints

- 1) *Context Viewpoint:*

a) *Users and Design Concerns:* The design features in this document are chosen to create value for several different kinds of users and stakeholders. These users interact with the software in different ways.

- 1) **Pilot:** Our intended user will be a pilot of a manned or unmanned aerial vehicle. This software is designed to provide increased and extensible functionality for onboard camera systems. Because our product is a proof of concept, there won't be a real pilot, but the pilot has been included as a user because the application of this design is to create better equipment for pilot use. This pilot will be 'using' our product and needs the video stream to meet the requirements outlined in the SRS.
- 2) **Distributor:** Rockwell Collins is the sponsor and primary stakeholder in this project. They will evaluate the product and decide where to use it. They will also decide the feasibility of using this design in production and whether continuing to use single board computers for real-time video processing is a good idea.
- 3) **Implementer:** The implementer has to set up the simple vision system on their given hardware. Our proof of concept should give some notion as to how to do this and also the feasibility and easiness of it. The concepts demonstrated in this design document can be used on different hardware, not just the NVIDIA Jetson.
- 4) **Algorithm Designer:** The algorithm designer interacts with our software in different ways than the end user and implementer. This user will create custom modules and link them together with our easy to use configuration system. This configuration should serve as a design for other systems.



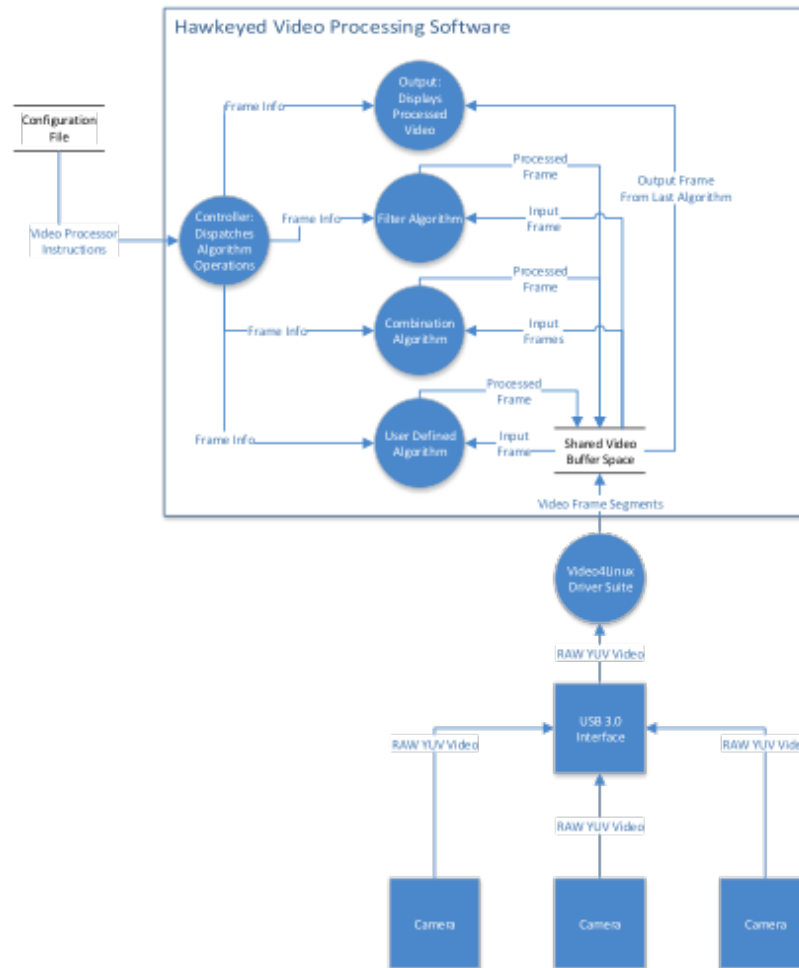
2) *Structure Viewpoint:* This section shows the high level organization of our software design.

a) *Design Concerns:* The primary design concern of this section is to show the interaction between critical components of our software design. This is valuable during implementation because it provides a reference for the required parts of the system, and also shows how the different pieces fit together.

b) Design Elements: There are three key design elements in our software design. These elements are important because they are the basis for the entire modular video processing system.

- 1) **Video Buffer:** Shared data space provided by video4linux or created internally. Video data is in YUV format and access is provided through a void pointer to a contiguous block of memory.
- 2) **Algorithm Module:** Algorithm modules are single purpose elements that operate on one or more input buffers and provide output to one or more output buffers. They consist of a C function which performs the operations, and a module definition structure which provides the name and expected input and outputs of a module.
- 3) **Algorithm Controller:** This element is responsible for ordering and executing Algorithm Modules. It is the primary routine in this software design. It reads from a JSON configuration file, creates a module execution tree, and manages the execution of concurrently running modules.

c) Structure Description: The HawkEye video processing software receives frames from the Video4Linux driver as a pointer to a buffer. The Algorithm Controller invokes each video processing algorithm in the order described by the configuration file, by passing them pointers to the buffers assigned to them. The video processing algorithms perform actions directly on the video buffer in order to avoid the performance hit from accessing additional memory. When the last video processing algorithm is complete, the buffer is flushed to the output device.



d) *Data Flow Diagram:*


3) *Interaction Viewpoint:*

a) *Design Concerns:* The main goal of this project is to determine how much video processing can be done on a SBC while maintaining a frame rate of 30 FPS and latency below 100ms. To this end, we as the user have two design concerns that need to be addressed.

- 1) **Latency measurement:** In order to keep overall latency under 100ms while maximizing the amount of processing being done, we need to know how long each algorithm takes to process a single frame, and also how long it takes from input to output.
- 2) **FPS Monitoring:** Our software must process video in real time. For the purposes of this project, that has been specified as 30 FPS by our SRS. This is separate from the latency requirement because our system may actually be operating on multiple frames simultaneously, so the frames per second may not necessarily dictate the total latency.

b) *Design Elements:* This viewpoint contains several design elements that depict major components of our software design.

- 1) **Camera Buffer:** Initially, camera buffer will store the data immediately loaded by the camera and will then be repurposed to hold the modified image later. This modified image will come from the custom algorithm code.
- 2) **Custom Algorithm:** The custom algorithm is loaded into the system when it starts up through the JSON configuration files and dynamically linked libraries. This portion of the system will perform operations on the camera buffer and log the time at start and time at completion. It will then write back to the camera buffer, the modified image.
- 3) **Log:** The log is a feed that will keep track of the frame rate and information about what processes are running. This should be viewable separately from the display.
- 4) **Display:** The last thing that should happen when we are rendering frames is the stream to the display. This will show the final, completely processed image. Our client has contacted us and told us that the display should only show the video stream, not other metrics like frames per second or operations per frame.
- 5) **Video4Linux Camera Driver:** This is our direct access point to the video stream from the USB camera. It enables us to receive video frames from any cameras supported by the Linux operating system.



c) *UML Sequence Diagram:*

4) *Information Viewpoint:*

a) *Design Concerns:*

- 1) **Configuration:** This viewpoint addresses how the user will configure our software. It shows how the modules are connected to each other and the information that needs to be provided for the software to create its internal structure.

- 2) **Internal Data Structure:** In order to implement the modular image processing system, our design needs an internal data structure in order to determine the execution path through the various modules.

b) Design Elements: There are two design elements in this viewpoint:

- 1) **Module Instances:** Module instances are C structures that describe either input device modules, output device modules, or algorithm execution modules. They consist of a name, type, and the names of the connected input and output modules.
- 2) **JSON Configuration File:** In order to create the internal data structure for our application, the user must write a configuration file in JSON format. The elements in this configuration file are detailed in an example below (Example JSON Configuration File).

c) Description and Rationale: The HawkEye video processing software needs to be able to handle various configurations of multiple camera inputs and produce multiple outputs, as well as performing multiple video processing algorithms. In order to do this, it needs to know what operations to perform in what order, and where to put the data. To do this, our software uses a tree-like data structure that connects the camera inputs, algorithm modules, and outputs. An input module definition specifies what modules a camera sends data to. Each of those modules then has its own output specification, which can either go to another algorithm module or an output module.

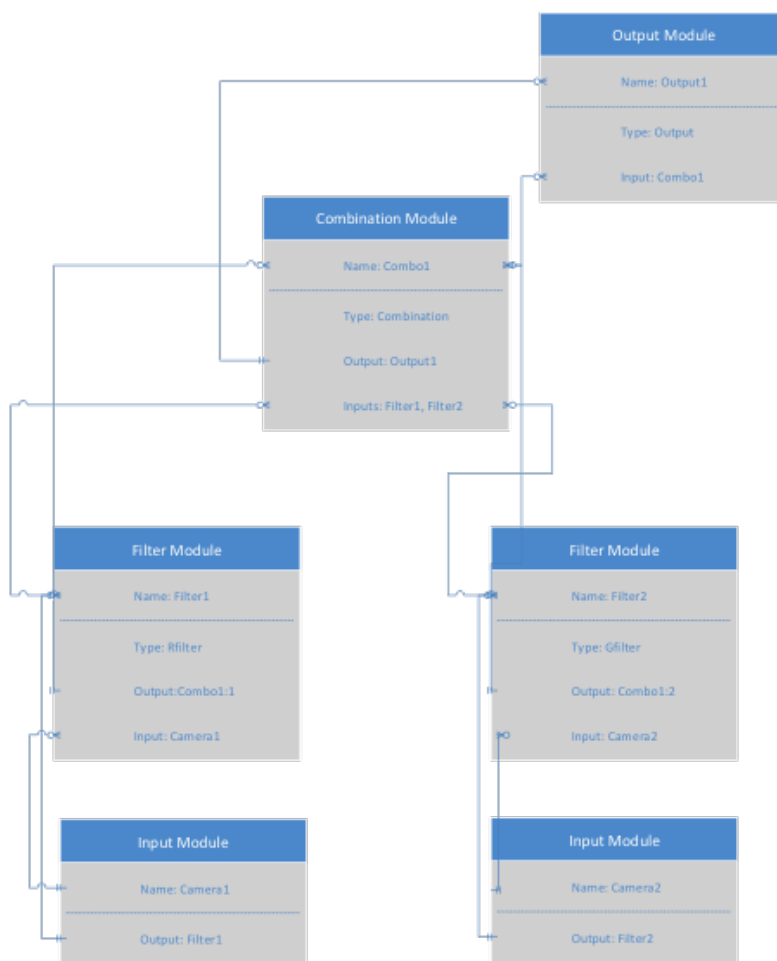
d) Sample Entity Relationship Diagram:

e) Example JSON Configuration File: The corresponding JSON configuration file for the above ER diagram.

```

1  {
2      "Output1":{
3          "type":"output" ,
4          "input": "Combo1"
5      },
6      "Combo1":{
7          "type":"Combination" ,
8          "inputs":" Filter1 , Filter2" ,
9          "output": "Output1"
10     },
11     "Filter1":{
12         "type":" Rfilter" ,
13         "input": "Camera1" ,
14         "output": "Combo1:1"
15     },
16     "Filter2":{
17         "type":" Gfilter" ,
18         "input": "Camera2" ,
19         "output": "Combo1:2"

```



```

20     },
21     "Camera1":{
22         "type":"Input",
23         "device":"/dev/video0",
24         "output":"Filter1"
25     },
26     "Camera2":{
27         "type":"Input",
28         "device":"/dev/video1",
29         "output":"Filter2"
30     },
31     (not used in this example but also an option)
32     "Tracker":{
33         "type":"Module",
34         "filename":"tracker.so"
35     }
36 }

```

The one rule for the JSON files is that each entity must be declared before its inputs; this top down approach enables the JSON parsing algorithm to generate the module tree in a single pass. The one exception to this rule is cyclical algorithms that operate on their own output, either directly or indirectly. Support for such algorithms may be implemented, but this is not part of the current project scope.

5) *State Dynamics Viewpoint:*

a) *Design Concerns:* Our program will be in multiple states while it is streaming and processing data from the camera to the display. In order to keep track of the states, we will have two state diagrams, one for the camera and a second for the software for processing.

b) *Design Elements:*

- 1) **Camera:** For the purposes of this viewpoint, Camera shall refer to the combination of the physical camera hardware and the system level driver that exists outside of the HawkEye video processing application.
- 2) **Software:** Refers to the HawkEye video processing application.

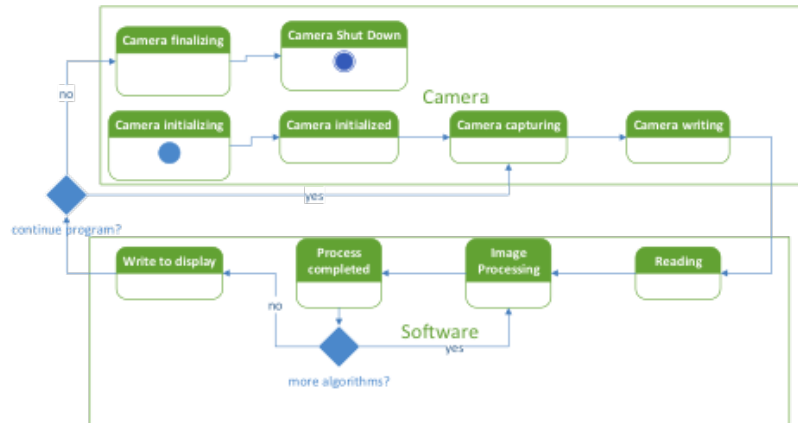
c) *Overview:* The camera must be initialized before we can capture frames from it. After initialization the camera is ready for image capture. During image capture, the camera will write to a buffer on the host system. Once a frame has been written to the buffer, the software can begin image processing using the first algorithm in its configuration. When the algorithm is completed, the software will check if there is another algorithm to be performed. These two states of image processing and checking for more algorithms will loop until all algorithms have been completed. We will then write the finished image buffer to the screen. After that, we will check if the program needs to continue running or if it has received a shutdown command. To shut down we will finalize the camera, deactivate it, and stay in that state.

d) *States:*

- 1) **Camera Initializing:** This is the starting state. Before we can capture frames from the camera, we must first initialize it. In this state we will also set up and synchronize the memory shared by the camera and our software.
- 2) **Camera Initialized:** After we have initialized the camera it will be ready for capture. The software will have to start the camera's frame capture.
- 3) **Camera Capturing:** The camera will take a moment to capture the data.
- 4) **Camera Writing:** The camera will have to write the data to memory.
- 5) **Image Reading:** Importing data from the camera's output into the software. This step should be simultaneous with camera writing as we are using the same memory.
- 6) **Image Processing:** The image data in the buffer will be processed in place by the modular video processing system.
- 7) **Process completed:** Once the algorithm is done operating, it will notify the controller that it has finished.

- 8) **Write to Display:** The software will write the finalized image buffer to the display, displaying the video.
- 9) **Camera Finalizing:** During this state, the camera is sending any remaining data and shutting down.
- 10) **Camera Shut Down:** The camera is turned off.

e) *State Transition Diagram:* Because the two state machines interact often, we have combined them in the same graphic. The following is our UML state machine:



V. DESIGN DOCUMENT CHANGES

A. Overall Description Changes

We ended up having a physical copy of the performance of the system at expo instead of working it into our software. This was because we were advised that using the software to measure it's own performance would lead to problems, leading us to measure the performance of the system using a high speed camera which can't be done automatically.

B. Specific Requirements

One of the goals of the original design document was to create requirements that were reasonably achievable. Because of this mindset, the requirements in the document did not change very much.

C. Information Viewpoint

We had originally intended to create a modular json system, but scrapped it later on because it wasn't very valuable and created unnecessary complexity.

VI. TECHNOLOGY REVIEW

A. Camera Interface (Hardware Level)

There are three different camera interfaces available on the Jetson development board. Each one offers a different set of benefits and challenges.

1) *USB 3.0:* USB 3.0 is a high-speed plug and play interface that offers a theoretical 384 MB/s data throughput. One USB 3.0 port is available on the Jetson, but additional ports can be added using the mini PCI Express expansion slot.

Benefits:

- Availability of USB 3.0 cameras
- Plug and play, with no camera configuration

Drawbacks:

- Jetson onboard USB 3.0 does not meet the USB 3.0 specification for bandwidth
- Some cameras do not have USB 3.0 drivers for ARM
- USB 3.0 only supports one camera per port due to bandwidth limits
- Increased cost due to board expansion

2) *Gigabit Ethernet:* The Jetson has one-gigabit Ethernet port available. Gigabit ports can handle 125 MB/s, however like USB 3.0 there are expansion cards available that add two additional ports.

Benefits:

- All gigabit cameras use a standardized interface and require no driver
- More high quality cameras available than the USB 3.0 offers
- Simple configuration with video4linux

Drawbacks:

- Most cameras require external power supplies
- Gigabit cameras typically cost more than the USB 3.0 alternatives

3) *MIPI Camera Interface:* The Tegra K1 processor on the Jetson has a direct camera interface, which is designed for onboard, low power camera modules. There are two MIPI interfaces, one 4 channel and one 1 channel interface, but any additional interfaces are unavailable.

Benefits:

- High speed, built in interface
- Low power consumption
- Lower latency because the video processing can be done before buffering

Drawbacks:

- Very limited number of cameras supported
- Specific to the Tegra K1, making it not portable
- Increased hardware complexity
- Increased software complexity due to lack of standardized drivers

4) **Decision:** We decided to use the USB 3.0 cameras because of the availability of low-cost high-quality cameras. Also, Rockwell Collins has several USB cameras they would like us to use. For our customer, it will be best to base the system around the USB 3.0 option. This may require adding a USB expansion card, but the other advantages of using USB 3.0 make it the first choice for camera connection.

B. Video Subsystem (Software Level)

We will have a framework for streaming the video from different components (camera, image processing unit, display).

1) *OpenCV:* We could use OpenCV to stream the video and setup processing of the video stream. This seems like the easiest approach, but would constrict us to the functions given by OpenCV, which might be enough for a proof-of-concept.

Benefits:

- Able to write scripts using python
- Can use OpenCL to take advantage of the GPU

Drawbacks:

- Hides functionality
- Questionable performance

2) *Direct Memory Handling:* We could handle the video stream ourselves on top of v4l provided by the Jetson, but would require some low-level knowledge of the Jetson's hardware. It would allow for more control over the video stream and give better performance.

Benefits:

- Fast performance
- Full control of operations

Drawbacks:

- Requires low-level knowledge of the Jetson architecture
- Complex, which make it much easier to develop bugs
- Not portable to other boards

3) *Gstreamer:* We could use Gstreamer, a free library with video functionality. Gstreamer is widely used as an industry standard, and works on different architectures. If we choose to use Gstreamer we are limited to the API it provides, which is pretty limited because it's meant for video conversion.

Benefits:

- Lightweight
- Portable
- Scriptable
- Easy pipeline-design

Drawbacks:

- Questionable performance
- May not be able to make use of hardware components

4) **Decision:** We will start developing with OpenCV, but if the performance is too slow we may switch to a different option. Video4Linux may be enough to easily move the stream to different components of the Jetson.

C. Image Processing

In order to demonstrate the processing power of the Jetson Tk1, we are going to implement some video processing algorithms.

1) *On-board GPU:* We could make use of the GPU on the Jetson to do image processing to test the performance. GPUs are commonly used for this type of processing and there are many ways to access the functionality of it.

Benefits:

- Easy to use
- Easy to maintain
- Easy to update

Drawbacks:

- Best performance (compared to the ISP)

2) *Integrate Stream Processor:* We could use the two ISPs on the Jetson to do image processing. ISPs are meant for video processing and would provide the best performance. It would allow us to bypass the memory if we connect the camera directly to the camera port. This would require tuning the video processing operations to use the ISPs interface.

Benefits:

- Fast
- Can bypass memory buffers

Drawbacks:

- Not-entirely portable
- Only one camera port

3) *PCIe GPU*: The Tegra has a PCIe port that could be used to install a more powerful GPU for video processing. This could be a good option if the Jetson doesn't have the required video processing power out of the box.

Benefits:

- Could potentially be very powerful

Drawbacks:

- Uses more power

4) ***Decision***: We are going to use the GPU for our image processing. Use of the GPU is widely supported by applications and is easy to code and maintain. Also OpenCV on the Tegra Tk1 directly supports it, and many other choices have support for GPU processing. It will also be able to handle our requirement of 30 frame rates per second processing.

D. Demo Interface

We will need a framework in order to display the video and user interface.

1) *QT Framework*: The QT framework is a stable library that provides user interface. It is cross-platform and has a simple API, but it is not a lightweight framework. Its ability to display video properly is undetermined.

Benefits:

- Easy to use
- Many features

Drawbacks:

- Bulky library
- Not specifically for video streaming

2) *OpenCV GUI*: OpenCV has various UI elements built in allowing for video display and user interaction. It is built into OpenCV and since we have decided to use OpenCV this makes it very simple.

Benefits:

- Part of the OpenCV library
- Easy to code

Drawbacks:

- Slow
- Limited API
- Limited functionality to control the look and feel

3) *OpenGL*: We could use OpenGL to display the video and interface. This would give us full control over the look and feel and it would be very responsive.

Benefits:

- Fast
- Portable
- Powerful (full control of display)

Drawbacks:

- Long development time
- Hard to maintain

4) **Decision**: We will use the OpenCV built-in GUI for the interface, as it is simple to use if we are already using it for the memory manipulation and video processing. It could even be used in combination with other choices, and would still be a good choice as the API is easy to use and is large enough to have all the features we want.

VII. TECHNOLOGY REVIEW DOCUMENT CHANGES

A. Technology Changes

1) *Camera Interface*: For our camera interface we ended up using the USB 3.0 camera setup. This made our code simple, but did require us to work with a USB 3.0 expansion card which had trouble supplying the necessary power to the cameras, requiring the camera's power supplies to be connected while in use.

2) *Video Subsystem*: For our video software subsystem, we ended up using a different system entirely: CUDA. We talked about OpenCV's performance issues in the Technology Review and they ended up constraining us too much. OpenCV was great for getting a simple program working, but it only ran at around 14 frames per second which was not viable. Writing lower level code in CUDA kernels and using the CUDA API allowed us to take more control over the memory in the TX1 resulting in higher frame rate.

3) *Image Processing*: We ended up using the on-board GPU for our image processing which is what we thought we were going to use. It was the best option because it was the simplest to code and investigate and it supposedly would meet the requirements, so we had no reason to believe it wouldn't be capable for the project and if it failed we wouldn't have wasted much time.

4) *Demo Interface*: In our technology review, we discussed a lot of software that we thought we might use to display the output of our camera system and decided to use OpenCV to create the project. We ended up using OpenGL instead. Our discussion of the benefits and drawbacks of OpenCV vs OpenGL in the technology review show why we eventually made the switch. The OpenCV GUI ended up being too slow compared to OpenGL. This was because of the benefit we list in OpenGL, saying that it was powerful. We had full control over the buffers that were sent to the window for rendering, which allowed us to employ memory mapping magic to reduce the number of memory copies and thus achieve the latency and framerate required for the project. While this was more expensive in terms of coding hours, it ended up being necessary for the project. Though attempting to use OpenCV was a worthwhile effort to learn more about our development environment quickly.

VIII. WEEKLY BLOG POSTS

During the duration of our project, our team posted weekly blog posts at the end of every week on Friday to recap progress, problems, and ideas we had come across during that week. We used these blog posts as a way to keep our TA and clients updated with the most current information about our project. This section contains all of our blog posts broken down by term, then by date of each Friday they were posted.

A. Fall Term 2015

1) October 16, 2105:

a) *First Meeting and Problem Statement Developed:* This week consisted of our initial meeting with our sponsors (we found out that we more or less have two of them), and we were also able to construct a complete problem statement of the project. The meeting occurred Monday in the the Valley Library. Our sponsors were able to drive down from Rockwell Collins in Wilsonville to have a face-to-face meeting. It was very nice to be able to have this meeting and to clarify many of our questions/concerns about the project. After the meeting was completed, the team split up different bullet points of the problem statement to work on. A Google document was created, where we were all able to collaborate and work on the document. Once the content was created we went through and edited it all to have the same voice and proper language. The final document was then created as a PDF and sent our sponsor to be signed. They were very happy with the outcome and feel we were able to capture an accurate and complete vision of the project. During this week we also created the group project page and blog site, and were able to do a little research involving the capabilities of the NVIDIA Jetson board. This concluded the week for our project.?

2) October 23, 2105:

a) *Requirements Document:* This week we drafted a requirements document and sent it off to out client for review. We are planning to meet with him soon over a WebEx VOIP, which I've never used. Hopefully that works out well. We're going to get a signature for our requirements document for next week. We also fixed some issues with SharePoint this week. We accidentally created two sites and in a very tense moment, deleted one. Also getting everyone the correct permissions on SharePoint has been difficult.

3) October 30, 2105:

a) *Requirements Document Continued...:* This week overall has been a slow one, waiting for review and input. It took until later this week to get some feedback from our client about the rough draft of the requirements document, which was needed for further advancement on the document. The feedback was very helpful in letting us know that we need to narrow our requirements down, and focus on additives once we can finish the main requirements. We have begun to make a more defined document with added descriptions and better defined requirements. Our document will be completed and signed by Wednesday, November 4th.

4) November 6, 2105:

a) Final Revision of Requirements Document: This week, due to an extension, we were supposed to turn in our final draft of the requirements document signed by our sponsor. On Tuesday, the class was informed that many of our requirements documents were not up to standard. We now have until Tuesday of this coming week to rewrite our documents and submit them again. We were given an IEEE standards document of how to correctly structure a requirements document. We took this document as a group and restructured our requirements. We added a cover page, a table of contents, and much more description of the project as a whole and each individual task as well. On Wednesday, we had our first meeting with our TA, Xinze. He was very helpful in giving us feedback for our original document. Once we had updated the old document into the newer format, we sent him a copy of it and got even more feedback. Once the document was completed, and the Gantt chart created according to our requirements was added, we sent it to our customer for his input. We are still currently waiting on a response. I believe that we have a pretty solid requirements document now compared to our original attempt. We also have several other assignments coming up in the few weeks, like a technology review, an elevator pitch, and a rough draft of our poster.

5) November 13, 2105:

a) Final Final Revision of the Requirements Document: We revised the requirements document yet again and submitted an almost finished version to our class. We then were able to contact our client for edits and are working on finally finalizing it. We also wrote up the tech document. It looks like we're going to start out using OpenCV and see if we can setup up the system using that and then optimize it from there. Also we might switch to using the TX1, a newer version of the TK1.

We're having a meeting with our client on Monday and should prepare for that.

6) November 20, 2105:

a) Rockwell Collins Meeting and Poster Development: On Monday this week, we met with Carlo and Weston from Rockwell Collins for our monthly check in. There were several items on the agenda:

- We went over the revisions we made to the Requirements document after receiving comments from Carlo via email last week
- We discussed alternative single board computers, including the Sapphire Tech Step Eagle board currently being used by Rockwell Collins for a similar project as well as the new Jetson TX1
- We formed a plan for evaluating camera options. Rockwell Collins is providing us with one USB 3.0 camera to test and see if we can get it to work with our Jetson. If it is not supported by the Jetson, we will begin looking at other options
- We discussed our Gantt Chart and received recommendations from Carlo about additional milestones we could put in to add more specificity to our project plan.
- Carlo and Weston are still unable to connect to our sharepoint site, so Hailey is working with them to see if she can get them set up.
- We discussed potential dates for our next meeting, which will be somewhere between December 14-16. Depending on if we get the hardware to work together, we may meet in person either in Portland or Corvallis.

After our meeting, we also created a rough draft of our project poster for the Engineering Expo.

7) *November 27, 2105:*

a) *Thanksgiving:* We had the second half of thanksgiving off and so we weren't able to accomplish much more than planning out a few meetings for the next week.

8) *December 4, 2105:*

a) *Design Document and Progress Report:* This week we created a design document and progress report. The design document followed the IEEE Std 1016-2009 format, which we went through and met on Monday to start brain storming. Our main concern with this document was identifying our viewpoints and creating corresponding diagrams for each. Below gives an overview of our viewpoints and the corresponding design languages:

- 1) **Context viewpoint:** This viewpoint shows the different potential users of the software and how they would interact with our system.
Design Languages: Use Case Diagram
- 2) **Structure viewpoint:** This viewpoint shows how the streaming video flows through the system, and identifies the internal and external data connection points in the system. This viewpoint also shows the components of the system which enable connection of cameras via USB as per the SRS.
Design Languages: Data Flow Diagram
- 3) **Interaction viewpoint:** This viewpoint shows the order of operations on processing a video frame, as well as how the timing is integrated to meet the performance benchmarking requirement.
Design Languages: UML Sequence Diagram
- 4) **Information viewpoint:** This viewpoint details how the modular video processing system determines the order of operations for the various elements and algorithms it can create. It shows how the system meets the requirement for modularity and the capability to use multiple input and output devices.
Design Languages: Entity Relationship Diagram
- 5) **State Dynamics Viewpoint:** While the algorithm implementations in our design are stateless, the system itself is not. This viewpoint shows the transitions between different states and provides a road map for different states which will need to be individually tested.
Design Languages: UML State Transition Diagram

The design document will be our roadmap to implementing our system to meet our requirements. We took special care in ensuring this document was done correctly as it is going to be a point of reference for the remainder of the year. We also created our progress report which covers a summary of the activities, problems, and solutions we came across during this Fall term. This winter break we will begin to familiarize ourselves with the Jetson development environment, and possibly connect our camera system to the board to find out how to connect the two environments.

B. *Winter Term 2016*

1) *January 8, 2016:*

a) HawkEye Crew is back at it: After a long, and much needed break, we are back into the grind of Winter term. This first week was busy for us. We began this week by finding that our Tx1 had a manufacturing problem, preventing it from connecting to wifi. We notified our instructor who gave us another board that turned out to be defective in the same way. Soon, we will get a Tx1 that definitely works correctly and start developing on it, but in the meantime, we've been working with the Tk1.

We installed graphics drivers on the Tk1 as well as drivers for the camera we currently have. We found that we need to recompile the Jetson Tk1's operating system in order to grab full 2048x2048 images from the camera, but we're able to grab images at a lower resolution using the FlyCapture API currently.

We got some extra hardware: a USB hub and two USB wifi dongles. Kevin supplied the hub and one of the USB wifi dongles.

We're trying to setup a meeting with Carlo this week to go over milestones and a timeline, which we built this week.

We also met with our TA today and learned what we need to do for the alpha stage. We must write up a document describing the process so far. The criteria for this document hasn't been defined clearly yet. We also have to make a video demonstrating our product. We were thinking of showing the camera in action and going over some of the code/interface for the project. This video has to explain each point of the requirements document which we should look over.

We still have to get a second camera working and setup a framework for image processing before the alpha stage.

2) January 15, 2016:

a) Beginning the Development Process: Week 2 of Winter term of is beginning to smooth itself out from last weeks discovery of the Jetson Tx1 manufacture bug and the difficulties of connecting to Wifi with the Tk1. We had to recompile the Tk1 with a special version of linux4tegra, called "Grinch", to get on output from the USB 3 camera. It proved to be not a simple task, as we still have some connection and display issues with the video. We began to realize that development on the Tk1 may not even be worth the struggle, and decided we need to get a working Tx1.

With Kevin to the rescue, we were able to pick up a new Tx1 from him that is fully functional, with all the bells whistles (including working on board Wifi). We were able to pick that up on Thursday, so we haven't had much time yet to try to get a working product between the Tx1 and our USB 3.0 camera, but plan to try and get that up and running before our meeting with Carlo on Monday via Webex.

Carlo requested a monthly milestone marker as to show what we should have to present to him and Weston at each of our monthly meetings. We created milestones Winter term, which will cover both our alpha and beta releases. Our rough idea of milestones is listed below.

January Meeting: January 11th

- Download graphics driver onto Jetson TK1
- Download device driver for camera onto TK1
- Display image onto monitor

February Meeting: Alpha Release ? February 11th

- Download all needed software onto Jetson TX1
- Display image from built-in camera on monitor
- Display image from RC camera and built-in camera
 - Stitch the two images together
- Have option for applying filters, but not actual apply the filters

March Meeting: Beta Release ? March 14th

- Fully functional filters
- Stretch Goals:
 - Modularity with JSON
 - Multiple monitor display
 - Object tracking
 - Display video operation metrics
 - Compare with different single board computers

??To finish off the week we all three sat down at meeting to revise our requirements document, as it is what we need to be moving forward with for the release of our alpha and beta product releases. Only a few things were changed, like making sure we have two working cameras for our system and to not specify only using the Tk1 or the Tx1. As we move forward we plan to rely more heavily on the Tx1, but want to allow ourselves the room to be able to fall back on the Tk1 if need be to complete our project. We will be uploading that document to the Sharepoint tonight and presenting the updated to our client on Monday at our monthly meeting.

3) *January 22, 2016:*

a) Meet-up and Decisions: We met with Carlo this week to show Carlo our current progress on the project. Our progress included a working video output from on board camera on the Jetson Tx1. We were able to display the image at 1080p with 30 or 60 fps. This video stream used embedded streamer code to transform the image buffer generated from the camera into a Video4Linux compatible format which then was read with OpenCV and displayed.

After the presentation of our project to our client, we went through the milestones we had sent Carlo last week. We are ahead of schedule with most of our February milestones complete. We just need to add a second camera and generate a few filters.

We have asked Carlo to send us the model number of the 4 other PointGrey cameras he has on hand. We also have asked for a list of typically used types of filters for aircraft vision systems.

We are to send Carlo a block diagram of the program flow from camera to display. This will be used to judge if it will be possible to incorporate RC's virtual camera system, easily, into our software.

We are also going to send some sample output of the video we have produced.

Lastly we will be researching a few options of cameras to add to our project and sending them to Carlo for his input on which ones would suite his image of the project.

We are focusing on getting two cameras' output displayed before moving on to fixing compatibility with other cameras or filters, and decided that using two USB 3.0 cameras would be the fastest way to yield this result.

4) January 29, 2016:

a) *PointGrey Camera on Point:* Last week we were able to get the ?camera that is onboard the Jetson Tx1 to display an image at 1080p and 30fps. We were having much more difficulty with the USB 3.0, PointGrey, camera that Rockwell Collins had supplied for us, but this week we managed to get it to work. ?

We were missing two big elements as to why we had so much trouble with getting a video output from it last week. We recompiled the Linux kernel to increase the USB buffer size to accommodate higher resolution image capture from the PointGrey camera. We also acquired code from PointGrey to interface between the proprietary camera API and openCV.

After we were able to get the video output working we implemented two basic edge detection functions using several of openCV's built-in functions to gain familiarity with the openCV API. Now that we have the USB 3.0 PointGrey camera working and running some filter capabilities, we have discussed the need for additional cameras that Carlo has available for us.

In class this week we went over the requirements needed for our Alpha release, which will contain a written document and 20-30 minute presentation of our progress with the project so far. The Alpha release must have all capabilities present, but do not need to have full functionality. For our project, this means having a video stream from multiple cameras displayed at decent frame rate and the ability to apply simple filters to the video stream. The next few weeks will comprise of writing the document needed for the release and moving forward with operating two cameras into one output.

5) February 5, 2016:

a) *Modularity and Multiple Cameras:* A major goal for this week was to get a second camera. This would allow us to be able to begin furthering our software to handle processing of multiple video streams into one output. Multiple cameras are a requirement that is needed for us to be able to have a complete project at Beta release, but something we wanted to be able to at least support with our Alpha release. We were able to get in contact with Carlo, and set-up a meeting allow us to (very generously) borrow two more PointGrey cameras with additional lenses as well. On Thursday Hailey drove up to Rockwell Collins in Willsonville ?to have an informal meeting/lunch with both Carlo and Weston to pick up the cameras.

Last weekend Ryan got a full modular system built within our software. ?Our software reads a JSON configuration file to set-up the modules??. This allows for more flexibility with the use of different cameras and their software. It also gives us the ability to different filter modules and only apply the ones we want a specific time. It could handle more flexibility with use of different types of cameras only being used at specific times or for specific reasons. For example, a long range camera may only be activated when needed to find the runway. It could then be turned on and processing when needed, but deactivated to not decrease the performance metrics when not needed.

At the beginning of the week, the modular system was able to output a static image which allowed us to demonstrate that the system was working. We then created a module to handle the PointGrey camera software. This allowed us to get two of the cameras up and running with one video output. We haven't gotten the USB PCIe card yet, so we can't actually implement dual cameras over USB3 yet, but were able to run one on USB 3.0 and the other on USB 2.0.? Even though one of the cameras was being ran with USB 2.0, we were still able to capture decent frame

rate and are below our required latency.? This weekend we will be fine tuning some of the video capabilities for our Alpha release presentation.

This week we also began the outline of our written midterm progress report and will hopefully have a solid rough draft for our meeting with Xinze on Monday.

6) *February 12, 2016:*

a) *Midterm Progress Report:* This week no real progress on the software was made because we were busy creating our Alpha release. As we had all of the required features implemented to have an alpha release we were able to focus on getting that out instead of rushing to finish any last features needed. The release consisted of a written progress report and a 20 min video presentation of the project to include a powerpoint and demonstration. The video presentation basically covered the topics within the written report and a demonstration of the output of our project. We used Prezi to create our power point presentation.

The report that covered a brief introduction of our senior project, our current progress, problems that have impeded our progress, and any preliminary results that have been gathered. It also described what is left to do before all of the requirements are met for the project along with the potential for some stretch goals if time allows. The report includes a few interesting pieces of code and a description of what this code does. It was constructed using the IEEEtran style guidelines.

On Thursday, we were also able to get our PCIe card that will allow us to hook up multiple USB 3.0 cameras. In the next weeks, this will be the direction we are moving with development. For our Beta release we are going to need to get two cameras running on USB 3.0, increase our performance metrics, and build our own filters.

7) *February 19, 2016:*

a) *Slow Week:* This week, we didn't accomplish much. We all had a lot of work outside the class this week. We met on Monday with our TA and explained to him about the progress we had made. He asked us about how our current alpha-level software works, all of the problems we had encountered, and our plan for completing the beta in the next few weeks. He also informed us about the coming class meeting next week. We are planning to meet this weekend to do more work on the beta project.

We also showed our client our video presentation and he approved of our progress.

8) *February 26, 2016:*

a) *Technical Stuff:* This week, we've been trying to make the visual frame rate more smooth on the Jetson. While we're capturing and processing the images, we have the required frame rate, but, display to the screen is very slow. Our metrics claim that we can get 30fps on the screen, but the resulting visual display is obviously not 30fps.

One solution we have to this problem is to use a low level API to display the images. We are planning to leave the capture and processing system in place, but simply replace the way we display the images, giving us more control over the display process. Our plan for replacement is to use GLUT which is a window system based on OpenGL. We have setup this replacement but haven't implemented the benefits of using GLUT which is zero copy. This means

writing the camera images directly to a buffer that is read by the GPU. because the GPU and the CPU share memory, this is possible.

Also we have very inconsistent results with these frame rates. Some of this was solved by writing more configuration details into our code, but the cameras sometimes remain in strange states and we are working on solving this.

We also found today that the Jetson TX1 isn't displaying on a monitor and wasn't responding to ethernet and had to reflash. This is mostly due to bad configuration on the last power down. Thankfully we were prepared for a situation like this and were able to reflash the TX1.

9) *March 4, 2016:*

a) FPS Up To 30 Using CUDA: This week we met our requirement for one camera by using CUDA libraries to process and handle the video buffers better. We currently have one camera outputting frames at 30 fps. We are trying to get this same performance with two cameras to fully meet our requirements.

We've also been getting the frame rate above 30 fps using more tricks in CUDA.

Another feature we added this week was a line detection filter, which can run without dropping our frame rate any significant amount. Eventually we will test the Jetson with these filters to see how many operations it will take to notice a drop in fps.

We also worked on the poster and end of term report a bit. We are having some trouble compiling tex files, but that should be sorted out soon.

We've also been having strange inconsistencies in the state of the cameras between operations of the software. This is mostly solved by unplugging and replugging the cameras, but we're trying to come up with a software solution to this problem.

10) *March 11, 2016:*

a) Everything's Working: This week we completed the requirements listed in our design document. We currently have 2 cameras streaming at 33 frames per second and 3 cameras running at 31 frames per second over USB3. We also got one camera running at around 100 frames per second but this program is unstable and crashes more than our program that runs at 30 frames per second.

We are planning to move forward using this program that grabs and displays frames at 100 fps, working out the bugs as we go along. This program can run at 100 fps because we spawn different threads for converting the color space from raw to RGB. Before, this was being done in the same thread as the capture thread which delayed the frame rate. Spawning new threads for processing had sped up this process.

We are planning to speed up frame capture even more by doing the color format conversion on the GPU using CUDA. This should get us to the maximum frame rate possible on the TX1.

C. Spring Term 2016

Development Timeline

- We chose the Jetson TX1 over the Jetson TK1 for development based on performance requirements
- We began with the on-board camera using GStreamer, but using multiple on-board cameras added unnecessary complications
- Started using OpenCV with USB3 cameras, but had low framerate due to poor memory provided by OpenCV
- The use of CUDA allowed us to fix our low framerate by sharing memory between the camera and GPU
- Replaced OpenCV's output with OpenGL
- Developed a multi-threaded program to further increase performance

Project Features

Hardware

- NVIDIA Jetson TX1 Single board computer with 256 core integrated GPU using less than 15 Watts
- 3 Point Grey high resolution cameras with various lenses and filters

Software

- CUDA GPU accelerated real time video processing modules
- Edge and feature detection algorithms
- Support for different camera modules including zoom, daytime, and low light cameras

**Rockwell
Collins**

Building trust every day



Image Processing Vision System for Manned and Unmanned Aircraft

Project Background

Rockwell Collins, our project sponsor, designs video systems to help pilots see better in tough weather conditions. These systems assist during flight operations, such as landing.

Some of these systems overlay images with graphics to help pilots locate important features like runways. Vision systems can also supply different functionality, such as applying NIR capabilities when in low light or even enabling a zoom lens to better identify a runway.

In order to provide pilots with this enhanced imaging, Rockwell Collins develops Field Programmable Gate Arrays (FPGAs) to perform image processing. This development process is very complex and costly. New vision enhancements can take weeks or months to develop using FPGAs. The FPGAs are currently Rockwell Collins' only option that uses low power and produces high quality video that pilots can use effectively.

Project Solution

Our goal is to provide a proof of concept for an alternative to Rockwell Collins' FPGAs. It must meet the performance metrics, provide faster implementation time, and reduce the cost of production.

We will test the capabilities of an NVIDIA single board computer, with our own image processing vision system. Our application will take in input from three USB 3.0 cameras, apply various image processing to the video stream, and display the results on a monitor.



TEAM MEMBERS

Hailey Palmiter
palmiterh@oregonstate.edu
Scott Griffy
griffys@oregonstate.edu
Ryan Kitchen
kitchenr@oregonstate.edu

PROJECT SPONSORS

Carlo Tiana
Weston LaBr

Results and Conclusions

Our proof of concept was able to meet the required performance metrics, have a faster implementation time, and reduce cost using the NVIDIA Jetson TX1.

Our performance metrics results are as follows:

- ~90 frames per second for three cameras
- 24 ms latency
- 1080p output resolution
- Video streams combined into one image output after GPU processing

Oregon State
UNIVERSITY

X. PROJECT DOCUMENTATION

XI. DISCOVERY OF NEW TECHNOLOGIES

XII. PERSONAL EXPERIENCES

A. Hailey Palmiter

B. Scott Griffy

C. Ryan Kitchen